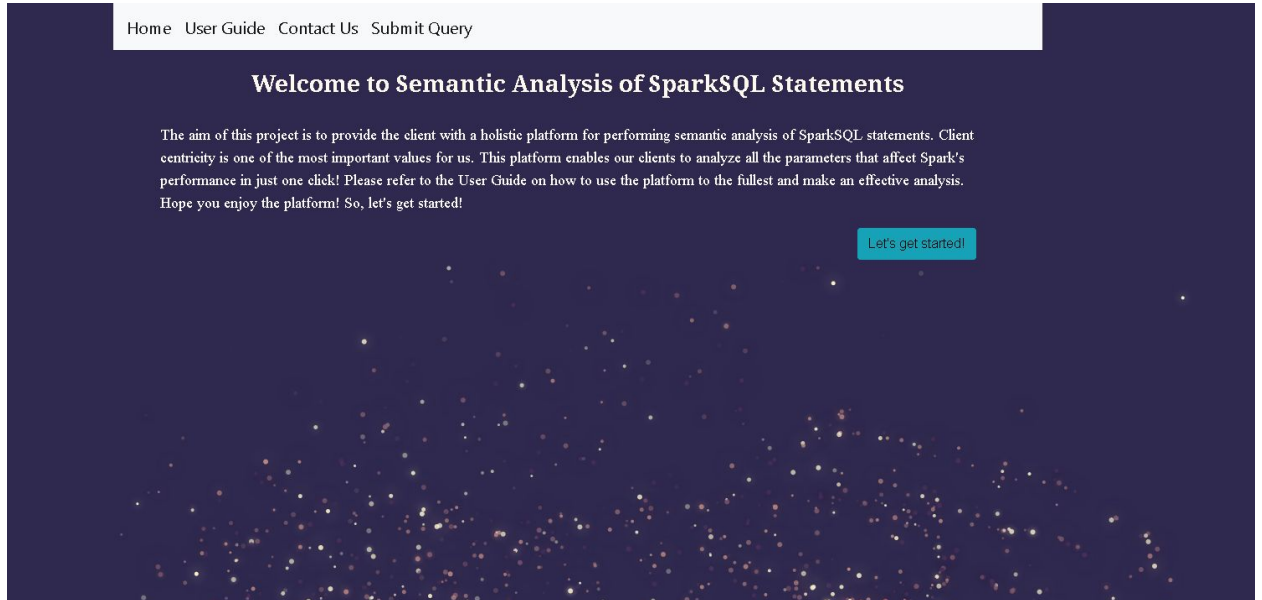


Semantic Analysis of Spark Sql Statements Project UI

1. Homepage



2. User Guide

[Home](#) [User Guide](#) [Contact Us](#) [Submit Query](#)

User Guide

This User Interface is designed to make the analysis of spark sql queries easy and convenient. Its a holistic platform which provides user with functionalities such as the query result, spark execution plans, the spark application UI and the hadoop cluster manager.

How to Use the application?

• Home Page

Click on 'Let's get started'

You will now be directed to the submit query page where you can enter your queries.

• Submit Query Page

We have provided 3 databases 'student', 'school', 'major' on which you can run your queries. You can view the databases by clicking on 'See All Databases'. Enter your query in the text box provided and click submit to run the query. If you want to rewrite the query click on reset. Now you will be redirected to the result page.

• Result page

Here you can see query that you have submitted.

1. Query Output

You can view the result of your query by clicking on the button.

2. Semantic Analysis

By clicking on this button you can view all 4 spark execution plans which are unresolved logical plan, resolved logical plan, optimized logical plan and physical plan.

3. Query Runtime

This will show the approximate runtime time of your query in milliseconds.

4. Spark UI

This button will open the Spark Application UI on a new tab. You can view the jobs, tasks and DAGs on the page which will help you for performance tuning.

5. Hadoop Application UI

You are now on the Hadoop Application page on a new tab which shows all the information about your hadoop cluster.

e.g. Active nodes, application running on cluster and metrics related to it.

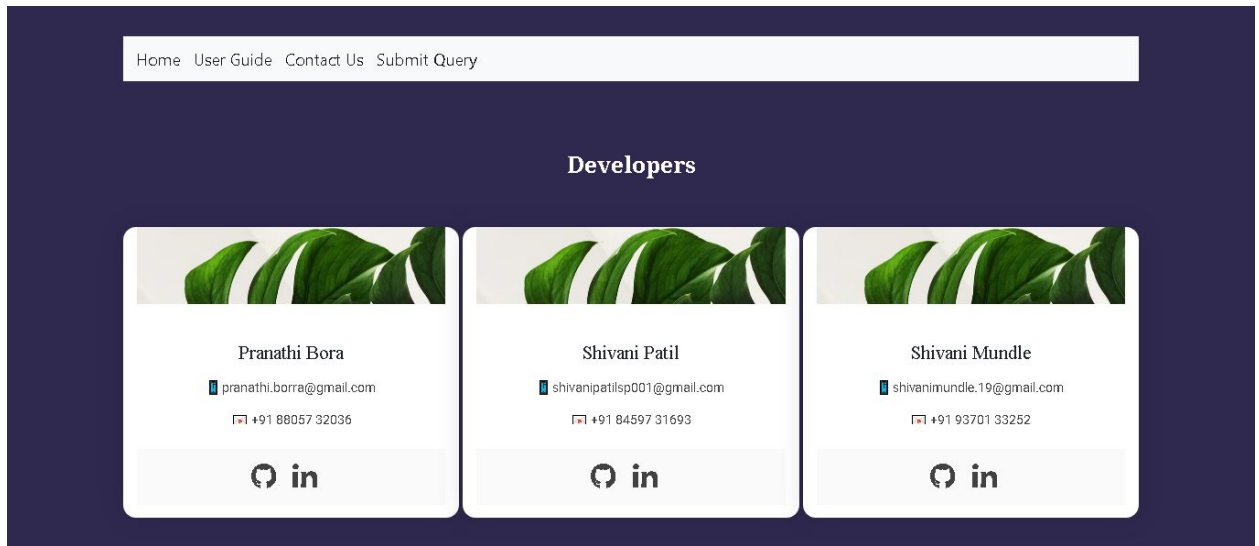
6. Hadoop NameNode UI

This will open the Hadoop NameNode UI where you can view information about the namenode and also see your directories and files on hdfs file system.

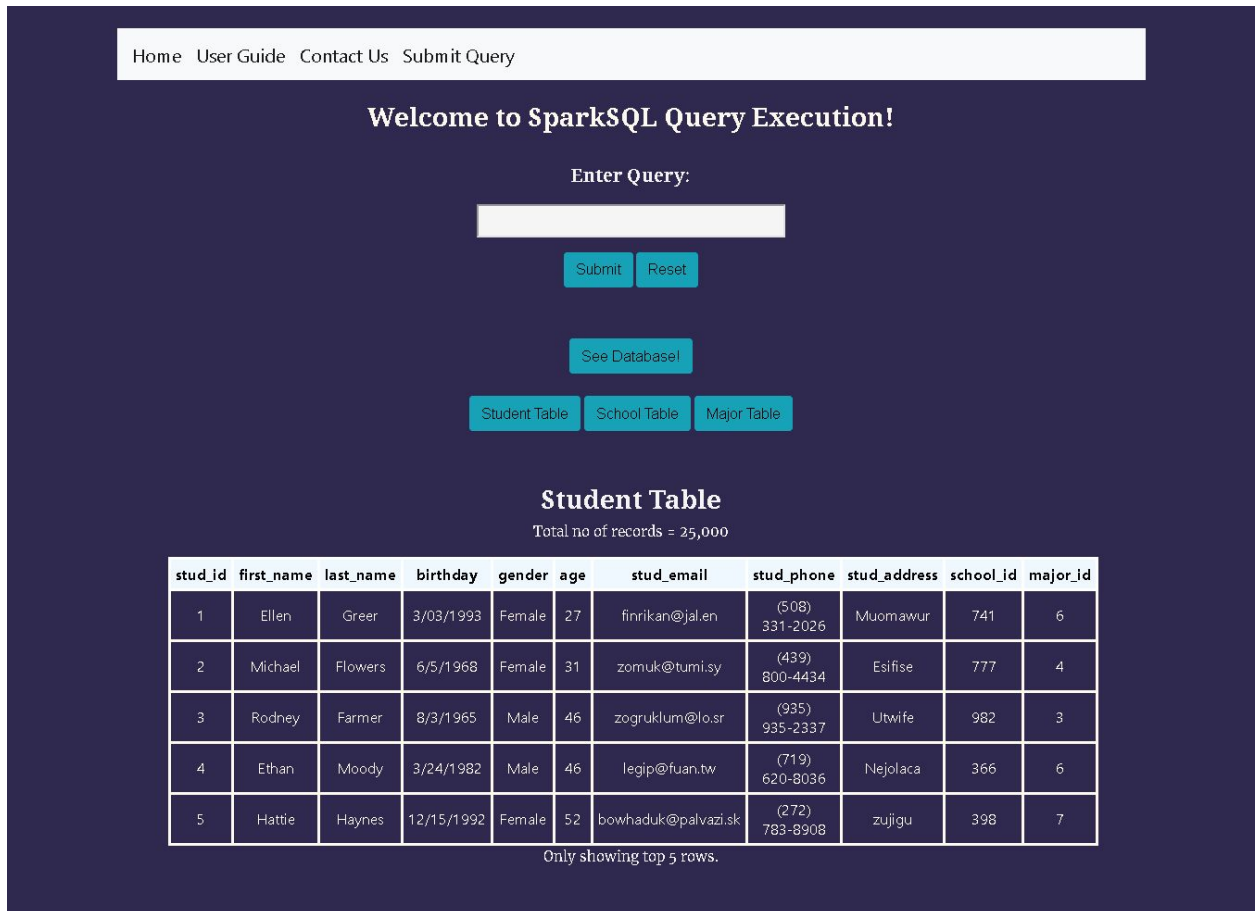
7. Submit Another Query

You will be redirected to the submit query page to run another query.

3. Contact Us Page



4. Analysis tool



5. Analysis

a. Submitted query

Query Results

Your Query: `select major_id,count(first_name) from student group by major_id`

[Query Output](#)[Semantic Analysis](#)[Optimization tips](#)[Run Time](#)[Spark UI](#)[Hadoop Application UI](#)[Hadoop Namenode UI](#)[Submit Another Query](#)

b. Query Output

Query Results

Your Query: `select major_id,count(first_name) from student group by major_id`

[Query Output](#)[Semantic Analysis](#)[Optimization tips](#)[Run Time](#)[Spark UI](#)[Hadoop Application UI](#)[Hadoop Namenode UI](#)[Submit Another Query](#)

major_id	count(first_name)
7	2518
3	2554
8	2485
0	2485
5	2503
6	2551
9	2502
1	2416
4	2537
2	2457

c. Execution Plans

1. Unresolved Logical Plan

Query Results

Your Query: select major_id,count(first_name) from student group by major_id

[Query Output](#)[Semantic Analysis](#)[Optimization tips](#)[Run Time](#)[Spark UI](#)[Hadoop Application UI](#)[Hadoop Namenode UI](#)[Submit Another Query](#)

[Parsed Logical Plan](#)[Analysed Logical Plan](#)[Optimized Logical Plan](#)[Physical Plan](#)

Unresolved Logical Plan

```
'Aggregate ['major_id'], ['major_id, unresolvedalias('count('first_name), None)]  
  
+- 'UnresolvedRelation `student`
```

2. Resolved Logical Plan

Query Results

Your Query: select major_id,count(first_name) from student group by major_id

[Query Output](#)[Semantic Analysis](#)[Optimization tips](#)[Run Time](#)[Spark UI](#)[Hadoop Application UI](#)[Hadoop Namenode UI](#)[Submit Another Query](#)

[Parsed Logical Plan](#)[Analysed Logical Plan](#)[Optimized Logical Plan](#)[Physical Plan](#)

Resolved Logical Plan

```
major_id: string, count(first_name): bigint  
  
Aggregate [major_id#67], [major_id#67, count(first_name#58) AS count(first_name)#109L]  
  
+- SubqueryAlias student  
  
+- Relation[stud_id#57,first_name#58,last_name#59,birthday#60,gender#61,age#62,stud_email#63,stud_phone#64,stud_address#65,school_#66]
```

3. Optimized Logical Plan

Query Results

Your Query: select major_id,count(first_name) from student group by major_id

[Query Output](#)[Semantic Analysis](#)[Optimization tips](#)[Run Time](#)[Spark UI](#)[Hadoop Application UI](#)[Hadoop Namenode UI](#)[Submit Another Query](#)

[Parsed Logical Plan](#)[Analysed Logical Plan](#)[Optimized Logical Plan](#)[Physical Plan](#)

Optimized Logical Plan

```
Aggregate [major_id#67], [major_id#67, count(first_name#58) AS count(first_name)#109L]
+- Project [first_name#58, major_id#67]
+- Relation[stud_id#57,first_name#58,last_name#59,birthday#60,gender#61,age#62,stud_email#63,stud_phone#64,stud_address#65,school_9]
```

4. Physical Plan

Query Results

Your Query: select major_id,count(first_name) from student group by major_id

[Query Output](#)[Semantic Analysis](#)[Optimization tips](#)[Run Time](#)[Spark UI](#)[Hadoop Application UI](#)[Hadoop Namenode UI](#)[Submit Another Query](#)

[Parsed Logical Plan](#)[Analysed Logical Plan](#)[Optimized Logical Plan](#)[Physical Plan](#)

Physical Plan

```
*HashAggregate(keys=[major_id#67], functions=[count(first_name#58)], output=[major_id#67, count(first_name)#109L])
+- Exchange hashpartitioning(major_id#67, 200)
+- *HashAggregate(keys=[major_id#67], functions=[partial_count(first_name#58)], output=[major_id#67, count#111L])
+- *Scan csv [first_name#58,major_id#67] Format: CSV, InputPaths: hdfs://localhost:9000/student_data.csv, PushedFilters: [], Re
```

d. Optimization tips

Query Results

Your Query: `select major_id,count(first_name) from student group by major_id`

[Query Output](#)[Semantic Analysis](#)[Optimization tips](#)[Run Time](#)[Spark UI](#)[Hadoop Application UI](#)[Hadoop Namenode UI](#)[Submit Another Query](#)

[Parsed Logical Plan](#)[Analysed Logical Plan](#)[Optimized Logical Plan](#)[Physical Plan](#)

Optimization tips

GroupBy

- GroupBy queries can be improved by tuning the number of partitions. It can lead to a significant performance improvement since it is a very expensive operation.
- The default number of partitions for aggregations is 200 in Spark. Try tuning this value to get an improved performance!
- If you have less distinct keys, then you can try repartitioning the data to less than 200 partitions. This can boost performance pretty much.
- You can repartition the data on a column value if you are planning on performing multiple complex aggregation functions and joins on the same column.
- You can also use repartition along with an expression to be performed like groupBy. The data gets shuffled according to the groupBy expression.
- If you don't want to use the repartition function, you can also set the SHUFFLE_PARTITIONS property to the desired number of partitions.
- GroupBy with RDDs is not optimized. If you're using groupBy with DataFrames you're pretty much good to go!

General

- Making simple changes to the system parameters might also improve the performance of SparkSQL statements!
- You can set spark.sql.codegen parameter to true as it compiles and creates java bytecode quickly for large queries but it may lag a bit on small queries.
- If you have small interim tables that are used repeatedly, caching them and then performing more complex operations can result in a significant performance improvement.
- Caching all the generated tables is not a good idea as cache memory is limited and may lead to the eviction of some blocks that are already in the cache.

e. Approximate run time

Query Results

Your Query: `select major_id,count(first_name) from student group by major_id`

[Query Output](#)[Semantic Analysis](#)[Optimization tips](#)[Run Time](#)[Spark UI](#)[Hadoop Application UI](#)[Hadoop Namenode UI](#)[Submit Another Query](#)

Approximate Run Time

169ms

f. Spark UI

 2.0.0

JobsStagesStorageEnvironmentExecutorsSQL

b4e2788b-0090-4392-8f7e-85820372... application UI


Spark Jobs (?)

User: Administrator
Total Uptime: 2.8 min
Scheduling Mode: FIFO
Completed Jobs: 14
[Event Timeline](#)

Completed Jobs (14)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
13	show at FirstController.java:99	2020/06/26 21:01:48	2 s	1/1 (1 skipped)	199/199 (1 skipped)
12	show at FirstController.java:99	2020/06/26 21:01:47	1 s	2/2	2/2
11	csv at FirstController.java:72	2020/06/26 21:01:45	36 ms	1/1	1/1
10	csv at FirstController.java:72	2020/06/26 21:01:45	38 ms	1/1	1/1
9	csv at FirstController.java:71	2020/06/26 21:01:45	38 ms	1/1	1/1
8	csv at FirstController.java:71	2020/06/26 21:01:44	37 ms	1/1	1/1
7	csv at FirstController.java:70	2020/06/26 21:01:44	35 ms	1/1	1/1
6	csv at FirstController.java:70	2020/06/26 21:01:44	39 ms	1/1	1/1
5	csv at FirstController.java:72	2020/06/26 21:01:22	0.1 s	1/1	1/1
4	csv at FirstController.java:72	2020/06/26 21:01:22	90 ms	1/1	1/1
3	csv at FirstController.java:71	2020/06/26 21:01:22	51 ms	1/1	1/1
2	csv at FirstController.java:71	2020/06/26 21:01:22	79 ms	1/1	1/1
1	csv at FirstController.java:70	2020/06/26 21:01:21	55 ms	1/1	1/1
0	csv at FirstController.java:70	2020/06/26 21:01:20	0.6 s	1/1	1/1

g. Hadoop application UI



All Applications

Logged in as: drwho

Cluster

[About](#)
[Nodes](#)
[Node Labels](#)
[Applications](#)
[NEW](#)
[NEW SAVING](#)
[SUBMITTED](#)
[ACCEPTED](#)
[RUNNING](#)
[FINISHED](#)
[FAILED](#)
[KILLED](#)
[Scheduler](#)

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes
0	0	0	0	0	0 B	8 GB	0 B	0	8	0	1	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[MEMORY]	<memory:1024, vCores:1>	<memory:8192, vCores:8>

Show 20 entries

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Tracking UI	Blacklisted Nodes
No data available in table											

Showing 0 to 0 of 0 entries

FirstPreviousNextLast

h. Namenode UI

Browse Directory

<div>/</div> <div>Go!</div>							
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	Administrator	supergroup	137 B	6/15/2020, 5:30:44 PM	2	128 MB	major.csv
-rw-r--r--	Administrator	supergroup	76.71 KB	6/15/2020, 5:29:18 PM	2	128 MB	school_data.csv
-rw-r--r--	Administrator	supergroup	2 MB	6/15/2020, 5:29:01 PM	2	128 MB	student_data.csv